



DREAMFORCE 2026 | ANNOUNCED SEPTEMBER 11, 2026

Job-Ready Agents

Seven Agentforce agents built for specific jobs — and how to tell when one fits better than an agent you build yourself.

THE SEVEN AT A GLANCE

Agent	The job it does	Availability
Casey help agent	Resolves customer service issues across voice, SMS, WhatsApp and web chat — FAQs, returns, account management and escalation to a person.	GA NOW
Paige IT & HR service agent	Resolves employee requests across Slack, portals and the tools employees already use.	GA NOW
Carter shopper agent	Helps shoppers discover and compare products, answers their questions, and converts with in-chat checkout.	GA NOW
Piper inbound pipe gen agent	Works across websites and inboxes to engage, qualify and convert inbound leads into pipeline for B2B sales and marketing.	GA NOW
Hunter outbound sales agent	Works a pipeline from research to outreach, collaborating with sellers over weeks and months on a single goal.	PILOT → NOV '26
Marshall supply chain agent	Orchestrates end-to-end back-office processes with deterministic execution and an audit record of every action.	GA NOW
Fin customer agent	Resolves complex customer-experience workflows across every channel, built on the Operator and Fin Apex models.	GA NOW

THREE THINGS TO KNOW BEFORE YOU COMPARE

1

They are Agentforce, not a separate product.
Same platform, same data, same security model — pre-configured for one job. You can rename them and extend them.

2

Job-ready is a starting point, not a ceiling.
Each ships with the skills, actions and data models its job needs, then gets tailored to how your company actually works.

3

Pre-built or custom is rarely the real question.
Most companies end up running both, with Multi-Agent Orchestration routing work between them as a job crosses roles.



What a “job-ready” agent actually is

The phrase causes more confusion than it needs to. A job-ready agent is not a new product line sitting next to Agentforce, and it is not a locked appliance. It is an Agentforce agent that Salesforce has already configured around a specific job — resolving a support case, qualifying an inbound lead, running a back-office process — so that the work of defining what the agent should know, what it is allowed to do, and how it should behave has largely been done for you.

Everything you could do with an agent you built yourself, you can still do with these. You are starting from a working configuration instead of an empty one.

What ships in the box

- **Skills and actions** for the job — the specific things the agent can do, already defined and wired to the right objects.
- **A data model** that maps to the Salesforce records the job depends on.
- **Channel connections** — voice, SMS, WhatsApp, web chat, Slack and portals, depending on the agent.
- **Behavior and guardrails**, including when to escalate to a person.
- **Grounding in Customer 360**, so the agent works from the customer context and business processes you already have in Salesforce.

What is still yours to own

- **Your data.** An agent is only as good as the knowledge, records and content it can reach. Gaps here are the most common reason a deployment stalls.
- **Your permissions model.** Every agent operates inside your business rules, permissions and security — which means those have to be right.
- **Your process.** Where the agent hands off, what it may do unsupervised, and what needs human approval.
- **Tailoring.** Tone, naming and edge cases. You can give an agent your own name so it reads as part of your brand and workforce.
- **Ongoing tuning.** Reviewing what the agent is getting wrong and improving it.

The part that is genuinely new: work that spans weeks

Most agents, including most of the ones companies built over the last two years, complete a task or a conversation and stop. Salesforce built a new long-horizon runtime for Agentforce so an agent can hold a goal across days and weeks instead. Hunter is the first agent to run on it; Salesforce has said more of the portfolio will follow, and that customers will be able to build long-horizon agents of their own.

Memory

Preserves context and progress across sessions, so the work does not reset when the interaction ends.

Durable execution

Keeps a plan running over time and lets the agent resume or course-correct as circumstances change.

Dynamic steering

Adapts the agent's behavior based on an individual user's feedback and direction.

What that looks like in practice. A seller asks Hunter to rescue at-risk deals before the end of the quarter. Hunter turns that objective into a measurable goal, builds a plan, and works toward it — deciding which tasks to complete, which tools and context it needs, and which guardrails determine when it can act on its own and when it needs the seller's approval. As things change, it adjusts the plan. The seller stays in control throughout.



The seven agents in detail

Three of the seven came into the portfolio through acquisition rather than being built from scratch, which is worth knowing: it means they arrive with a track record outside Salesforce rather than as first-generation software.

Agent	What it does	Where it runs	Origin	Status
Casey Help agent	Resolves customer service issues end to end, with pre-built support for FAQs, returns, account management and escalation to a person.	Voice, SMS, WhatsApp, web chat	Salesforce	GA NOW
Paige IT & HR service agent	Resolves employee IT, HR and administrative requests without routing them to a queue.	Slack, employee portals, internal tools	Salesforce	GA NOW
Carter Shopper agent	Guides shoppers through discovery and comparison, answers product questions and completes the purchase in the conversation.	Commerce storefronts, chat	Salesforce	GA NOW
Piper Inbound pipe gen agent	Engages website visitors and inbound leads, qualifies them and converts the ones worth converting into pipeline.	Website, inbound email	Qualified	GA NOW
Hunter Outbound sales agent	Researches accounts, runs outreach and stays with a seller on a goal across weeks or months. First agent on the long-horizon runtime.	Email, CRM, seller workflow	Salesforce	PILOT → NOV '26
Marshall Supply chain agent	Orchestrates back-office and supply chain processes across systems, people and other agents, with deterministic execution and a full audit record.	Back-office systems, multi-step workflows	Regrello	GA NOW
Fin Customer agent	Handles complex customer-experience workflows, built on Operator and the Fin Apex models trained specifically for customer experience.	Every customer channel	Intercom	GA NOW

A note on what “GA now” does and does not mean

Generally available means you can buy it and switch it on. It does not mean the agent will be good at your version of the job on day one. Every one of these grounds itself in your data and runs inside your permissions, so the quality of the output tracks the quality of what it can reach. Companies that get value quickly almost always did the knowledge and data work first — or picked a first use case narrow enough that the data was already clean.

Hunter is the exception on timing: it is in pilot, with general availability announced for November 2026. Salesforce's own published materials note that any unreleased capability should not be the basis for a purchase decision.



How these agents actually work

The common objection to any packaged AI agent is that it is a chatbot with a name on it. The distinction that matters is what sits between the model and your business, because that is what determines whether an agent can be trusted with a real customer or a real ledger.

Four things that make an agent an agent

Grounding	Agents are connected to Customer 360 and work from the customer context and business processes you already have in Salesforce, rather than from a model's general knowledge. Answers come from your records, your knowledge base and your approved content.
Permissions	Every agent operates within your business rules, permissions and security. An agent cannot see or do something a user in that context could not see or do — which is why your permissions model is a prerequisite, not an afterthought.
Determinism	Agent Script, Salesforce's open-source language for agent behavior, lets you combine AI reasoning with fixed rules. You decide which parts of a process the model gets to reason about and which parts always execute the same way — the difference between a helpful suggestion and a financial transaction.
Audit	Actions are recorded. Marshall in particular provides an audit record of every action it takes, which is the baseline expectation for back-office and regulated processes.

The governance layer around them

At Dreamforce, Salesforce described the six capabilities it believes any enterprise needs in order to run agents safely at scale — the Trusted Enterprise AI Harness. It is a useful checklist whether or not you buy any of it, because it names the things that have to be true before an agent touches revenue or customers.

Trusted models

Choose the right model per use case, including Koa, Salesforce's CRM-tuned reasoning model — currently in pilot — alongside frontier models.

Trusted context

Data, metadata, semantics, knowledge, real-time signals and memory — plus an understanding of how work actually gets done.

Trusted agency

The plans, steps and runtime logic that turn a desired outcome into an executable path.

Trusted actions

Secure connections to the applications, APIs, workflows and processes the agent needs in order to finish the job.

Trusted governance

Data quality, classification and policy enforcement, so only appropriate information is ever exposed to AI.

Trusted security

Identity, permissions, privacy, data protection and runtime security around what an agent may access and do.

And how they get better

- **Multi-Agent Orchestration** routes work across agents so they operate as one team when a job crosses roles, systems or stages of the customer journey. *Generally available.*
- **Agent Optimizer** helps teams build and refine agents, test performance and analyze session traces to find what to improve. *Generally available October 2026.*
- **AI Skills in Agentforce Coworker** let an employee teach the agent how to complete a task once, then scale that know-how across the workforce. *In pilot; generally available October 2026.*



“Why wouldn’t we just build it ourselves?”

It is the right question, and the honest answer is that you can. If you have Agentforce, you already have everything these agents are made of. Nothing here is withheld from the builder and sold back to you.

So the decision is not *packaged software versus platform*. It is a narrower and more practical trade: for this particular job, is the configuration work worth doing yourself? For a process that is unique to your company, it usually is — that process is part of why customers choose you. For resolving a password reset or answering “where is my order,” it usually is not, because there is no version of that job that is distinctively yours.

Which way each situation points

Start with a job-ready agent	Build your own	Expect to do both
• The job is one that most companies do roughly the same way. • You want something running this quarter, not next year. • You do not have an AI or platform team with spare capacity. • The job spans channels you would otherwise have to wire up individually. • You want a reference implementation to learn from before you build.	• The process is proprietary, and doing it your way is a competitive advantage. • The work depends on systems or data models unique to your business. • No agent in the portfolio maps to the role you actually need filled. • You have the team, and the ongoing tuning capacity, to own it properly. • Regulatory or contractual constraints dictate a specific implementation.	• A job-ready agent covers the common 70% and a custom agent handles your exceptions. • You start packaged to prove value, then extend the same agent rather than replacing it. • Different functions move at different speeds, so service goes packaged while operations builds. • Multi-Agent Orchestration routes work between them as a job crosses roles.

What the trade actually costs you either way

	Job-ready agent	Agent you build
Time to something working	Days to weeks. Published customer deployments run from under an hour for a first internal agent to about six weeks for a full shopper journey.	Weeks to quarters, depending on scope, team availability and how much of the data and permissions work is already done.
Who does the work	Configuration and tailoring, typically an admin plus a business owner.	Design, build, evaluation and ongoing tuning — a standing commitment, not a project.
Fit to your process	Good for standard work, adaptable at the edges. Poor if your process is genuinely unusual.	Exact, by definition. You are building to your own specification.
Who owns quality	Shared. Salesforce maintains the agent; you own the data, permissions and process it depends on.	You, entirely — including the evaluation harness and the regression testing nobody budgets for.
What it costs	Consumption or per-user licensing, visible from day one. Scales with use.	Build cost plus the same platform consumption. The build cost is the part usually underestimated.



A way to decide, one job at a time

Score the specific job you are considering — not your company, and not AI in general. The same organization will score differently for customer service than it does for supply chain, which is exactly the point.

For this job, how true is each statement?	Not true 0	Somewhat 1	Very true 2
Most companies in our industry handle this job in broadly the same way.			
How we do this job is not something customers choose us for.			
We need something in production within the next one or two quarters.			
We do not have an AI or platform team with capacity to own a new build.			
The knowledge and records this job depends on are already in reasonable shape.			
The job runs across more than one channel or system.			
We would rather learn from a working implementation than design from a blank page.			

10–14

Start with a job-ready agent

This job is standard, the clock matters and you are not differentiated by how you do it. Take the packaged agent, tailor it, and spend the team you saved on something that is differentiating.

5–9

Start packaged, plan to extend

A job-ready agent will cover the common cases and you will outgrow parts of it. That is a normal path, not a failure — extend the same agent rather than starting over.

0–4

Build it yourself

This job is specific to you, or the portfolio has no agent for the role. Use the platform directly, and borrow the patterns from a job-ready agent where they apply.

Two scores that should override the total

- **If the data question scored zero, fix that first.** An agent grounded in incomplete or contradictory knowledge will produce incomplete or contradictory answers, and it will do so at volume. This is the single most common reason a pilot stalls, and no amount of packaging compensates for it. It is also not a reason to wait indefinitely — scope the first use case to an area where the data is already good.
- **If the differentiation question scored two, be careful about building anyway.** Teams frequently decide a process is unique to them when what is actually unique is the accumulated workaround. Worth testing that assumption before committing a year of engineering to it.



What results look like, and how fast

The figures below are the ones Salesforce has published by name. They are real deployments rather than modeled projections, which also means they reflect those companies' data, processes and scope — not a promise about yours.

50%**Engine**

of chat inquiries fully resolved by its help agent, named Eva

60%**Perk**

of sales pipeline built by its outbound sales agent, Hunter

70%**Autism Queensland**

of administrative requests resolved by Paige

90%**Hibbett**

of core shopper journeys handled by Hibbett AI, live in six weeks

4×**Asana**

the conversation volume via Piper, deployed in 45 days on average

79%**Anthropic**

of the conversations Fin sees are resolved autonomously

Two deployments worth looking at in detail

Fulton Bank — speed of first value

One hour from first hearing about Agentforce to having an agent set up. Power users were on it within a week. Six weeks later, 3,000 employees were using it, across more than twenty use cases. Salesforce has since published 80,000 hours saved and complaint resolution times cut by four days. The instructive part is the sequence, not the speed: one narrow use case first, then width.

Southwest Airlines — what extending looks like

Began with a focused assistant answering customer questions, freeing reps for the interactions that needed empathy. Then expanded into name changes, receipt requests and travel vouchers, orchestrated as subagents inside one unified experience. Salesforce reports a seven-fold return on investment and \$6 million in projected annual savings. Their stated method was to watch real conversations and adjust rapidly rather than plan in advance.

The pattern underneath all of them

Salesforce describes a three-step sequence, and the published deployments follow it closely: **pick one use case and deploy it fast**; **customize and extend** that use case across the business; then **observe and optimize** in real time. The companies that moved quickly did not begin with a transformation program. They began with one narrow job where the data was already good enough, then widened.

The honest caveat. Every figure on this page comes from a company that had already done the underlying work — their knowledge base was in order, their permissions were coherent, and someone owned the outcome. A resolution rate is a result, not a starting condition. The useful question to ask about any of these numbers is not “can we get that?” but “what did they have in place before they could?”



How this is priced

There is no separate price list for job-ready agents. They are priced through the same models as the rest of Agentforce, and most organizations end up on more than one at once – consumption for customer-facing work, per-user licensing for employees. Everything below is from Salesforce’s published pricing page as of September 29, 2026.

Model	Published price	What it covers	Best suited to
Salesforce Foundations	\$0	Agentforce Builder, Prompt Builder, Agent Script, Coworker and Vibes.	Getting started and proving a use case.
Flex Credits	\$500 per 100,000 credits	Customer-facing and employee-facing agents, Voice, Slackbot and Digital Wallet. A standard action is 20 credits; a Voice action is 30.	Scaling across mixed use cases where you want cost tied to actual work done.
Conversations	<i>Quoted</i>	Customer-facing agents and Digital Wallet. Flat per-conversation pricing rather than per action. Salesforce does not publish the rate.	Customer-facing volume where predictable flat pricing is preferred.
Agentforce add-on	\$125 user / month	Unmetered Agentforce usage for employees, the full AI suite in the flow of work, AI analytics and Prompt Builder. For Sales, Service, Field Service and Operations.	Employee-facing agents at predictable cost.
Industries add-on	\$150 user / month	Everything in the Sales and Service add-ons, plus industry-specific AI. For Industries Clouds.	Regulated and industry-specific deployments.
Agentforce 1 Editions	from \$550 user / month	The add-on included, plus 2.5M Flex Credits per org per year.	Organizations standardizing on agents across functions.
Agentforce User License	\$5 user / month	Access for every employee with metered usage and limited CRM object access. Requires Flex Credits.	Broad, light-touch access across a whole workforce.

What consumption actually looks like

Salesforce publishes worked examples. These are illustrative and exclude other consumption such as Data 360 credits, but they give a realistic sense of order of magnitude.

Use case	Actions	Cost each	Assumed volume	Per month
Customer self-service – “where is my order”	2	\$0.20	20 requests / day	\$120
Service – case management in Salesforce	3	\$0.30	100 users, 3 cases / day, 20 days	\$1,800
IT support – new employee onboarding questions	1	\$0.10	20 new starters, 5 questions each	\$10

The question to ask before you sign anything. Consumption pricing means the bill tracks success: an agent that resolves more work costs more. That is the right incentive, but it needs modeling rather than estimating. Ask for your own volumes to be run through the rate card, ask which buying model fits – Pre-Purchase, Pre-Commit or PayGo – and make sure Digital Wallet is switched on from day one so consumption is visible before it is a surprise.



The questions people are actually asking

Isn't this just a chatbot with a name on it?

The name is branding; what sits underneath is not. These agents ground in your Customer 360 data, run inside your permissions, can execute deterministic steps via Agent Script, and log what they did. A chatbot answers. These complete work and leave a record of it.

We already built agents in Agentforce. Does this replace them?

No. It is the same platform, so what you built keeps working. Multi-Agent Orchestration is designed for exactly this situation — routing work between agents, however each one was created, when a job crosses roles or systems.

Our data is not clean enough. Should we wait?

Do not wait for the whole estate; do not skip it either. Pick a first use case where the knowledge is already in reasonable shape and prove it there. The data work is real, but it goes faster when it is scoped to one job with a visible result attached.

What happens when an agent gets it wrong?

That is what the guardrails and escalation paths are for, and why the deterministic controls matter. Decide up front which steps the model may reason about and which always execute the same way, and where the agent must hand to a person. Then watch real conversations and adjust — that is what the companies with the best published numbers did.

How is this priced, and what if usage spikes?

Consumption, per-user licensing, or both. See the previous page. Digital Wallet gives near-real-time visibility with threshold alerts, which is the control for exactly this worry. Model your own volumes rather than accepting a benchmark.

Can an agent work outside Salesforce?

Yes. Paige runs in Slack and employee portals; Casey covers voice, SMS and WhatsApp; Piper works across websites and inboxes. Salesforce also opened platform capabilities to external AI interfaces, so an agent does not require a Salesforce UI to do its job.

How long until we see something?

For a narrow internal use case on clean data, days. For a full customer-facing journey, weeks — Hibbett went live on 90% of core shopper journeys in six weeks, and Piper deployments average 45 days. The variable is rarely the software; it is how ready the data, permissions and process owner are.

Sources

Salesforce Newsroom, "Salesforce Expands Agentforce With a New Portfolio of AI Agents Built for High-Value Work," September 11, 2026. • Salesforce Newsroom, "Five Big Ideas Coming Out of Dreamforce 2026," September 24, 2026. • Salesforce, Agentforce Pricing, accessed September 29, 2026.

Please note. This brief summarizes publicly announced information and is provided for discussion purposes only. Pricing, packaging, features and availability are subject to change; confirm current terms with your Salesforce account executive before making any decision. Hunter is in pilot at the time of writing, with general availability announced for November 2026 — purchase decisions should be based on capabilities that are currently available. Customer results cited are those companies' published outcomes and are not a representation of the results any other organization will achieve. Salesforce, Agentforce and Slack are trademarks of Salesforce, Inc. Other marks are the property of their respective owners.